

STRATEGIC ALIGNMENT & REALIZATION ANALYSIS

Modernizing Defense Innovation Proposal vs. DoDI 5000.02, 5000.85, and 5000.88

Prepared for Program Management Review | April 2026

Executive Summary

This analysis maps and evaluates the alignment between the Defense Innovation Design Proposal and three governing DoD acquisition instructions: DoDI 5000.02 (Adaptive Acquisition Framework), DoDI 5000.85 (Major Capability Acquisition), and DoDI 5000.88 (Engineering of Defense Systems) — herein referred to collectively as MDAP policy. The assessment is drawn from two independent analytical reviews and the proposal itself.

The proposal's central contribution is not compliance theater — it is a set of specific engineering mechanisms (SBD, RDSG, PSM) that answer each policy mandate with a concrete operational method, not a documentation artifact.

The finding is unambiguous: the three-pillar architecture achieves primary realization of the most operationally significant requirements across all three policy instructions. Where coverage is partial — chiefly in workforce readiness and IP strategy under DoDI 5000.88 — the gaps are structural rather than conceptual, and are addressable in the Systems Engineering Plan.

Strongest coverage: DoDI 5000.02 and 5000.85, where SBD and RDSG together address two of the most persistent failure drivers in Major Defense Acquisition Programs — the failure to effectively complete system info-data definition and realization, and the absence of machine-enforceable requirements governance.

Most nuanced coverage: DoDI 5000.88, where the proposal's SysML v2/KerML infrastructure is a genuine architectural contribution, but its role is structural enabler rather than first-class compliance mechanism — the distinction a Milestone B reviewer would press on.

1. Analytical Methodology

The alignment assessment uses a two-axis evaluation framework. The first axis measures whether the proposal provides a primary realization (direct operational mechanism fulfilling the mandate) or a

supporting realization (structural enabler that makes compliance possible but does not constitute the mechanism itself). The second axis identifies coverage gaps — areas where the mandate is acknowledged but the proposal's response is thin or requires supplemental documentation.

Three source documents were reviewed: the primary Design Proposal, a focused regulatory alignment review, and a strategic cross-reference mapping prepared against MITRE SE best practices, INCOSE Handbook provisions, and SysML v2 implementation guidance. Findings were synthesized against the actual regulatory text of the three policy instructions rather than summary characterizations.

The resulting Alignment Matrix (Section 3) grades each proposal pillar against each applicable policy instruction. The DoDI 5000.88 Deep-Dive Table (Section 4) resolves the individual section-level provisions where the proposal's engineering architecture has specific, traceable applicability.

2. Proposal Overview

The proposal rests on three interdependent pillars. Understanding their interaction is essential to reading the alignment correctly — these are not modular compliance checkboxes but an integrated engineering approach where each pillar amplifies the others.

Set-Based Design (SBD)

SBD responds to a well-documented failure pattern in defense acquisition: programs that commit to a fixed solution concept early in development — selecting a single alternative and progressively firming it through successive milestone decisions — leave themselves with little recourse when that concept encounters unforeseen technical risk or evolving operational requirements. Rather than selecting a single solution at the outset, SBD maintains a bounded set of viable design alternatives simultaneously and employs Multiobjective Value of Information (MOVIO) analysis to determine when — and only when — it is epistemically justified to narrow that set toward a committed solution.

It is worth acknowledging a point of analytical nuance here. GAO reports and independent program assessments consistently show that cost growth and schedule slippage in MDAPs are rooted not simply in premature commitment to a solution concept, but more fundamentally in the failure to effectively complete system info-data definition and realization — that is, the requirements specifications, architecture descriptions, design descriptions, drawings, and interface control documents that constitute the authoritative technical baseline. SBD does not resolve that execution problem on its own. What it does provide is a disciplined framework for determining when a program has accumulated sufficient knowledge to make those commitments responsibly, and a governance mechanism that prevents organizational pressure or schedule momentum from forcing commitment before that threshold is reached.

Requirements-Driven Semantic Gateway (RDSG)

The RDSG transforms prose requirements from passive documentation artifacts into active governance infrastructure. Using SHACL (Shapes Constraint Language) and SWRL (Semantic Web Rule Language), requirements become machine-readable and machine-enforceable — violations surface automatically without human review cycles. Technical Performance Measures (TPMs) are tracked continuously against acceptance thresholds, and Value-Focused Thinking (VFT) metrics score system state against stakeholder objectives in real time.

The SysML v2/KerML layer beneath the RDSG is architecturally significant. The shift from XMI file exchange to a federated REST API means that multi-vendor teams query a single authoritative repository rather than maintaining potentially divergent local copies. This is the operational definition of the 'digital authoritative source of truth' that DoDI 5000.88 mandates but rarely sees executed at this fidelity.

Predictive Scoring Metrics (PSM)

PSM introduces mathematically grounded reliability metrics for AI and autonomous components — specifically the Brier Score for probabilistic calibration and the Mean Time Between Hallucinations (MTBH) metric for language-model reliability. The MTBH frames AI hallucination risk as an engineering reliability problem amenable to standard reliability engineering methods: measurement, threshold-setting, and circuit-breaker intervention when the threshold is violated.

This framing is consequential for program management. MTBH gives Program Managers a metric they can brief at a milestone review — quantified, bounded, and auditable — rather than a qualitative assurance that the AI was 'sufficiently tested.'

3. Alignment Matrix: Proposal vs. MDAP Policy

The matrix below maps each of the three proposal pillars against each MDAP policy instruction, indicating whether the alignment constitutes a primary realization (direct mechanism fulfilling the mandate) or a supporting realization (structural enabler). Empty cells indicate the pillar does not apply to that policy instruction's provisions.

DoD Standard	SBDSet-Based Design	RDSGReq.-Driven Semantic Gateway	PSMPredictive Scoring Metrics
DoDI 5000.02 Adaptive Acquisition Framework	PRIMARY Deliberate decision delay preserves pathway flexibility; no full program reset on technology insertion	PRIMARY Closed-loop requirements evolution uses AI behavioral failures as continuous refinement signals	—
DoDI 5000.85	PRIMARY	PRIMARY	—

DoD Standard	SBDSet-Based Design	RDSGReq.-Driven Semantic Gateway	PSMPredictive Scoring Metrics
Major Capability Acquisition	MOVOL analysis defers irrevocable design commitments, directly attacking the root cause of the Nunn-McCurdy breach	Quantitative trade-off evaluation across cost, schedule, performance, and risk before EMD entry	
DoDI 5000.88 Engineering of Defense Systems	SUPPORTING SysML v2/KerML federated repository replaces brittle XMI exchange; enables machine-verifiable baselines	PRIMARY SHACL + SWRL automated governance operationalizes the digital authoritative source of truth mandate	PRIMARY Brier Score and MTBH provide mathematically grounded AI/autonomy assurance per §3.6.e

Legend

PRIMARY	Direct mechanism fulfilling the mandate — the proposal provides the specific how
SUPPORTING	Structural enabler for compliance — indirect but essential foundation
—	Not applicable to this standard/pillar combination

4. Policy-by-Policy Analysis

The three policy instructions examined here — DoDI 5000.02, 5000.85, and 5000.88 — represent the governing framework for Major Defense Acquisition Programs (MDAPs) and are the appropriate reference set if MDAPs are the primary target audience. It is worth noting, as a collaborator has suggested, that broadening applicability to the full range of defense acquisition pathways — including the Software Acquisition pathway, Middle Tier of Acquisition (MTA), and Urgent Capability Acquisition (UCA) — would require additional alignment analysis against the policy instruments governing those pathways. That is a productive avenue for future work and would substantially expand the proposal's addressable program population.

4.1 DoDI 5000.02 — Adaptive Acquisition Framework

DoDI 5000.02 establishes the Adaptive Acquisition Framework (AAF) to deliver effective, affordable solutions promptly. Its operative emphasis is on speed of delivery, continuous adaptation, and frequent modular upgrades — explicitly as a corrective to the rigid milestone structures inherited from Cold War procurement culture.

SBD as the primary compliance mechanism

The proposal's most direct response to AAF is SBD's 'deliberate decision delay' — the systematic deferral of irrevocable design commitments until epistemic uncertainty is resolved. This is not procrastination but a formal engineering discipline: MOVOL analysis quantifies when additional

information would change the design decision, and decisions are deferred until that threshold is passed.

For AAF's modular upgrade requirement, SBD's maintained design set provides the natural insertion point. A new AI model or updated sensor specification enters the set and is evaluated against existing alternatives using the same MOVOL framework. No program reset is required because no single design has been irrevocably committed. This is precisely the 'continuous adaptation' the framework mandates.

RDSG provides a continuous feedback loop

DoDI 5000.02 requires iterative software development with continuous feedback loops. The RDSG operationalizes this through a closed-loop requirements evolution architecture in which AI behavioral failures are treated as structured information signals that propagate back into the requirements baseline. The loop is not a human review cycle — it is automated through the SHACL/SWRL governance layer, meaning that system behavior continuously tests requirements validity and surfaces violations in real time.

Coverage assessment: Full primary realization across both applicable pillars. PSM has no direct applicability to AAF provisions — its scope is AI/autonomy risk management, which falls under DoDI 5000.88.

4.2 DoDI 5000.85 — Major Capability Acquisition

DoDI 5000.85 governs Major Defense Acquisition Programs (MDAPs), with particular focus on cost, schedule, and performance trade-off management during the Technology Maturation and Risk Reduction (TMRR) phase and formal milestone decisions. Its most consequential enforcement mechanism is the Nunn-McCurdy statute, which triggers mandatory program restructuring or termination when cost growth exceeds defined thresholds.

SBD and the Nunn-McCurdy challenge

Nunn-McCurdy breaches are among the most visible indicators of MDAP failure, triggering mandatory program review and restructuring when cost growth exceeds statutory thresholds. The root causes documented in GAO oversight reports point consistently to breakdowns in system info-data definition and realization — the failure to generate and maintain authoritative requirements specifications, architecture descriptions, design descriptions, drawings, and interface documentation in a form that supports sound engineering decisions. This is not merely a matter of committing to a solution concept too early; it is often the failure to execute that definitional work effectively at all.

SBD's contribution in this context is not a cure for execution failure, but a governance discipline that addresses one specific and preventable contributor: organizational and schedule pressure that forces formal commitment to a fixed solution concept before the program has accumulated sufficient knowledge to make that commitment responsibly. When a program formally selects a single alternative

during TMRR, and that selection subsequently encounters unforeseen technical risk or revised operational requirements, the cost of reverting at EMD-stage maturity is severe. MOVOI analysis provides a quantitative criterion for when commitment is epistemically justified, reducing the likelihood of that specific failure mode without claiming to resolve the broader execution challenges that GAO reports document.

The MOVOI framework provides the formal decision criterion. It evaluates alternative designs simultaneously across cost, schedule, performance, and risk objectives, generating a quantitative answer to the question: Is there enough information to commit? If not, the set is preserved, and uncertainty reduction activities are prioritized. This is risk management executed at the architectural level rather than applied retrospectively at milestone reviews.

RDSG as pre-EMD trade-off infrastructure

DoDI 5000.85 requires programs to evaluate alternatives and reduce technical risk before entering Engineering and Manufacturing Development. The RDSG supports this by maintaining a live, queryable requirements baseline against which alternative solution concepts can be evaluated at any time without manual documentation overhead. Trade studies are not batch events preceding a milestone — they are continuous, enabled by the semantic governance infrastructure.

Coverage assessment: Full primary realization. The alignment here is not only regulatory — it represents a genuine engineering solution to the most expensive failure mode in MDAP history.

4.3 DoDI 5000.88 — Engineering of Defense Systems

DoDI 5000.88 is the most technically demanding of the three policy instructions. It establishes systems engineering policy across the full acquisition lifecycle, mandating digital engineering ecosystems, technical baselines, Independent Technical Risk Assessments (ITRAs), and specific provisions for AI and autonomous systems. The proposal's alignment here is strong but differentiated — PSM achieves primary realization of the AI/autonomy provisions; RDSG achieves primary realization of the digital engineering mandate; SBD's role is structural enabler rather than direct compliance mechanism.

RDSG and the digital authoritative source of truth

The DoDI 5000.88 digital engineering mandate is often satisfied in form but not substance — organizations produce a model-based artifact that is structurally a digitized Word document. The RDSG substantively fulfills the mandate by making requirements machine-readable, machine-enforceable, and continuously tested against system behavior. The SHACL/SWRL governance layer means that the authoritative source of truth is not declared by fiat but maintained by automated verification.

The SysML v2/KerML transition compounds this. KerML's 'occurrence semantics' — modeling system elements as entities with full life histories suitable for direct simulation — makes the digital baseline not

just a configuration record but a simulation substrate. This supports the PDR/CDR evidence generation that DoDI 5000.88 Section 3.5.a requires.

PSM and the AI/autonomy mandate

DoDI 5000.88 Section 3.6.e explicitly requires systems engineering to address hazards and risks specific to autonomy, artificial intelligence, and unmanned systems. Most programs satisfy this with qualitative assurance documentation. The proposal answers it with two quantitative metrics: the Brier Score (probabilistic calibration of AI agent uncertainty representation) and MTBH (reliability-theoretic framing of hallucination risk).

The Circuit Breaker pattern operationalizes MTBH: AI-generated engineering claims are validated against the SysML v2 repository before propagating into the authoritative baseline. The circuit breaker trips when the MTBH threshold is violated, preventing unverified AI output from corrupting the digital twin. This is systems engineering applied to AI reliability, not just software testing.

Coverage gap: workforce readiness and IP strategy

The proposal's treatment of DOD 8140 workforce alignment and intellectual property strategy — while present — reads as checklist compliance rather than integrated engineering design. The staffing mapping to IASAE levels is stated but not detailed, and the IP strategy's commitment to preventing vendor lock-in is asserted without a mechanism. Both are addressable in the Systems Engineering Plan but represent genuine gaps in the proposal's current framing.

Coverage assessment: Primary realization of the two most technically significant provisions (digital engineering, AI/autonomy risk). Supporting realization of the SysML v2/KerML baseline. Partial coverage of workforce and IP provisions — addressable at SEP stage.

5. DoDI 5000.88 Deep-Dive Cross-Reference

The table below provides section-level traceability between proposal elements and DoDI 5000.88 provisions, drawing on the MITRE Systems Engineering Guide and INCOSE Handbook alignment documented in the strategic cross-reference review.

Proposal Element	DoDI 5000.88 §	Strategic Significance	Alignment
Mission Engineering	§3.3	Establishes mission baseline and Concept Design Trade Matrices before material development, reducing requirements creep at MDA approval	Primary
Technical Baseline (SysML v2/KerML)	§3.4.b	Federated REST API replaces brittle XMI; machine-verifiable requirement closure rather than documentary assertion	Primary
MOSA Compliance	§3.7.a	Satisfies 10 U.S.C. 2446a-c on severable components; enables incremental upgrades and cross-Service reuse	Primary
ITRA Integration	§3.5.b	Provides independent data for Milestone B certification, focusing on critical technology and manufacturing maturity	Supporting
STPA / System Safety	§3.6.e	Addresses emergent unsafe system-level outcomes from software interactions — beyond component-level failure analysis	Primary
RDSG (SHACL + SWRL)	§3.4 (Digital Auth.)	Operationalizes the digital authoritative source of truth mandate; machine-enforceable requirements governance without human intervention	Primary
MTBH / Brier Score	§3.6.e (AI/Autonomy)	Mathematically grounded AI reliability metrics allow program managers to brief on hallucination risk at milestone reviews	Primary
HiL Co-Simulation (FMI)	§3.5.a (PDR/CDR)	Acausal-causal model integration generates Certification Evidence satisfying ITRA criteria for production readiness	Primary

6. Key Findings and Recommendations

6.1 Findings

- The proposal achieves full primary realization of DoDI 5000.02 and 5000.85 requirements through SBD and RDSG. These two pillars address the most consequential MDAP failure drivers with specific, quantitatively grounded mechanisms rather than process documentation. SBD's contribution is disciplined knowledge-based commitment; RDSG's is machine-enforceable system info-data governance.
- DoDI 5000.88 alignment is strong but differentiated. RDSG's SHACL/SWRL governance and PSM's MTBH/Brier Score metrics constitute primary realizations of the digital engineering and

AI/autonomy mandates, respectively. SBD's role in the 5000.88 context is as a structural enabler through the SysML v2/KerML infrastructure layer.

- The MTBH metric is a genuine contribution to the field. Framing AI hallucination risk as a reliability engineering problem — with a defined metric, threshold, and circuit-breaker intervention — provides Program Managers with something auditable and defensible at milestone reviews.
- The SysML v2 transition from XMI to federated REST API is not cosmetic. It is the infrastructure decision that determines whether the digital authoritative source of truth holds up across multi-vendor program teams or collapses into siloed governance.
- Coverage gaps exist in workforce alignment (DOD 8140) and IP strategy. These are structural rather than conceptual gaps and should be resolved in the Systems Engineering Plan before Milestone A.

6.2 Recommendations

- Milestone A preparation: Develop explicit staffing-to-IASAE-level mapping as part of the SEP workforce section. The current framing is adequate for a proposal but insufficient for IRT scrutiny.
- IP strategy: Specify the mechanism by which the Government retains data rights to SysML v2 repository artifacts. Assertion is insufficient; the strategy should name the contract vehicle and data rights category.
- MTBH threshold calibration: Define system-specific MTBH thresholds before EMD entry. The metric framework is sound; the program-specific calibration is the remaining work.
- STPA integration: The proposal references STPA for system safety but does not detail how STPA loss scenarios integrate with the RDSG's continuous requirements validation. This integration should be made explicit in the SEP's safety section.
- Acquisition pathway expansion: The current analysis is scoped to MDAP policy (DoDI 5000.02, 5000.85, 5000.88). A follow-on alignment pass against the Software Acquisition pathway, Middle Tier of Acquisition (MTA), and Urgent Capability Acquisition (UCA) instruments would materially broaden the proposal's applicability and funding addressability.
- IRT readiness: The three-phase implementation roadmap (Digital Authoritative Source of Truth → Virtual Integrated System → Physical Anchor) maps directly onto ITRA criteria. Document this mapping explicitly to accelerate IRT preparation for Milestone B.

7. Conclusion

The Defense Innovation Design Proposal is an unusually well-grounded regulatory document. Its three-pillar architecture does not merely describe alignment with MDAP policy — it provides the specific operational mechanisms that each policy instruction's mandate implies but does not specify. Where the policy defines what must be achieved, the proposal provides the how.

The alignment is strongest where the engineering problems are hardest: providing a knowledge-based commitment framework to reduce the risk of forcing fixed solution selections prematurely, replacing document-based system info-data governance with machine-enforceable semantic infrastructure, and quantifying AI reliability risk in terms that program management can use at milestone reviews. These are not incremental improvements to existing practice. They represent a substantively different approach to defense acquisition engineering.

The identified gaps — workforce readiness depth, IP strategy mechanism, STPA-RDSG integration — are real but proportionate. They are the expected delta between a proposal document and a fully developed Systems Engineering Plan, not evidence of conceptual weakness. Addressed in the SEP, they do not impair the proposal's readiness for Milestone A consideration.

The proposal is ready for Milestone A consideration pending SEP development that addresses the four identified gaps: DOD 8140 staffing detail, IP rights mechanism, explicit STPA-RDSG integration, and acquisition pathway expansion analysis.